

# QueryTorque

## VLDB Competitive Positioning & Game Plan

*Reasoning-First, Cross-Engine SQL Optimization*

### Paper Thesis (One Sentence)

"Existing AI query optimizers rely on iterative search (LITHE), static retrieval (R-Bot), or physical plan steering (LLM-QO). QueryTorque is the first system that **reasons from first principles** with a **structured knowledge framework**, and generalizes across database engines without retraining or retrieval."

Two pillars. One clean contribution. No scope creep.

### Pillar 1: Reasoning-First (No RL)

The architectural claim: SQL rewrites should be **reasoned**, not **searched**. QueryTorque produces a correct rewrite before it touches the database — no trial-and-error, no feedback loops, no dependence on unreliable cost signals.

#### Why this matters to reviewers

- **LITHE's loop is implicit RL with a broken reward signal.** Cost estimators can be off by orders of magnitude with skewed data. Optimizing against unreliable cost signals is not reasoning.
- **LLM-QO depends on planner feedback.** Without hint injection support (e.g., DuckDB, Snowflake), plan steering is not applicable. QueryTorque's reasoning is self-contained.
- **Token and time efficiency.** If LITHE takes 15 rounds to converge, QueryTorque takes 1 reasoning pass. Quantify this: tokens consumed, wall-clock overhead, DB round-trips.
- **Interpretability.** When reasoning fails, you can explain why. When an RL loop converges on a bad rewrite, nobody knows why. Trustworthiness matters in production DB systems.

### Pillar 2: Cross-Engine Generalization

The experimental claim: QueryTorque works across engines **without retraining, corpus swapping, or hint injection**. Every competitor is over-fitted to Postgres.

#### The DuckDB Experiment

**Run R-Bot against DuckDB (TPC-DS 1TB).** Its retrieval corpus is full of Postgres-specific advice (e.g., `SET enable_nestloop=off`) which does not apply to DuckDB and may degrade performance. LLM-QO's hint injection is not supported in DuckDB.

| Engine            | QueryTorque  | R-Bot       | Outcome                  |
|-------------------|--------------|-------------|--------------------------|
| Postgres (TPC-DS) | ~15x speedup | ~5x speedup | QueryTorque wins 3x over |
| DuckDB (TPC-DS)   | ~20x speedup | Fail / ~1x  | R-Bot cannot generalize  |

**Narrative:** "Existing methods are engine-overfitted. QueryTorque uses logical reasoning with a portable knowledge framework, making it the first cross-engine AI optimizer."

**VLDB reviewers value generalization over raw speed.** Winning on both is the ideal position, but if forced to choose, generalization is the stronger narrative.

## The Knowledge Framework (Curated, Not Auto-Generated)

Frame the knowledge base as a **structured expert rule library**. Describe the format (Gap Profiles), show examples, explain how a DBA extends it for a new engine. This gives reviewers reproducibility, correctness guarantees, and a clear extensibility story — without any hallucination concerns.

| Aspect                 | R-Bot (Static Retrieval)         | QueryTorque (Curated Framework)       |
|------------------------|----------------------------------|---------------------------------------|
| <b>Source</b>          | Scraped from manuals/forums      | Expert-authored, validated rules      |
| <b>Correctness</b>     | Unverified (forum advice varies) | Human-validated, deterministic        |
| <b>Portability</b>     | Engine-locked (Postgres corpus)  | New engine = new rule set (2hrs work) |
| <b>Extensibility</b>   | Requires new corpus per engine   | Structured format, DBA-extensible     |
| <b>Reproducibility</b> | Depends on retrieval quality     | Format published, rules inspectable   |

**Key point:** R-Bot got into VLDB '24 with scraped forum advice. A curated, validated knowledge framework is strictly better. The contribution is the reasoning architecture and how it interfaces with knowledge — not where the knowledge comes from.

**IP protection:** Publish the format and examples. The specific rules are operational advantage, not withheld science. No reviewer can object to this — it's like a trained model vs. training data.

## Benchmark Strategy: DSB Over TPC-DS

DSB (Decision Support Benchmark) is the right choice. Here's how to frame it:

### Why TPC-DS understates your advantage

Standard TPC-DS assumes column independence (no correlations). The Postgres optimizer handles this reasonably well, making it hard to show dramatic wins. The competition looks better than it actually is on TPC-DS.

### Why DSB reveals the real gap

DSB introduces data skew and correlations (e.g., customers in 'CA' buy more 'Winter Coats'). This is where standard optimizers — and systems that trust them — break down:

- Cost estimators assume uniform distribution and produce wildly wrong estimates on skewed data.
- LITHE's iterative loop optimizes against these inaccurate estimates, compounding the error.
- R-Bot's static rules don't account for data-dependent patterns.
- QueryTorque's reasoning can spot the skew and choose rewrites accordingly.

**Action:** Highlight results on **DSB queries 21, 36, and 78** (notorious skew queries). Briefly explain the skew pattern for each so reviewers see it's principled, not cherry-picked.

## Competitor Landscape

For reference: each competitor's approach, limitations, and benchmark claims.

| System            | Approach   | Limitation   | Their Benchmark   |
|-------------------|--|--|---|
| LITHE EDBT '26    | <b>Trial-and-error loop.</b> Generates rewrites, checks syntax, checks optimizer cost, iterates. | Wastes tokens in generate-check-fix cycles. Trusts the DB cost estimator as reward signal — but cost estimates are unreliable, especially with skewed data. No reasoning about <i>why</i> a rewrite works. | <b>13.2x GM speedup</b> (on self-selected "Hefty" queries only) |
| R-Bot VLDB '24    | <b>RAG retrieval.</b> Fetches static rules from manuals and forums, applies them.                | Cannot generate new optimization logic — limited to what exists in the corpus. Corpus is engine-specific: Postgres tuning advice does not transfer to DuckDB and may degrade performance.                  | <b>~5x speedup</b> (Postgres-only evaluation)                   |
| LLM-QO SIGMOD '25 | <b>Plan steering.</b> Forces the DB to use specific join orders via hint injection.              | Operates at the physical plan level — steers join orders via hint injection rather than rewriting SQL. Not applicable on engines without hint support (DuckDB, Snowflake, BigQuery).                       | <b>68% latency reduction</b> (DSB on Postgres only)             |

## Positioning Matrix

How QueryTorque differs across every evaluation dimension. Each row is a potential reviewer question.

| Dimension          | LITHE             | R-Bot            | LLM-QO             | QueryTorque                  |
|--------------------|-------------------|------------------|--------------------|------------------------------|
| Reasoning          | RL-style loop     | None (retrieval) | Planner feedback   | <b>First-class reasoning</b> |
| Engine Portability | Single engine     | Corpus-locked    | Hint-dependent     | <b>Cross-engine</b>          |
| Optimization Level | Syntactic rewrite | Rule application | Physical plan only | <b>Logical SQL rewrite</b>   |
| Knowledge Source   | None              | Static manuals   | None               | <b>Curated expert rules</b>  |
| Correctness        | Post-hoc check    | Unvalidated      | Execution-based    | <b>Validated before exec</b> |

## Anticipated Reviewer Questions

1. ***"What if reasoning gets it wrong? At least RL self-corrects."***

Answer: When reasoning fails, the failure is interpretable and diagnosable. When an iterative loop converges on a suboptimal rewrite guided by inaccurate cost estimates, the failure is silent. We validate rewrites for semantic equivalence before execution.

2. ***"How much overhead does the reasoning add?"***

Answer: Report wall-clock time, tokens per query, and DB round-trips. Pre-empt this by including an overhead analysis table in the paper.

3. ***"Why not combine plan steering (LLM-QO) with your SQL rewriting?"***

Answer: They're complementary. Our rewriting is engine-portable and operates at the logical level. Plan steering is engine-specific. A combined system is interesting future work, but our contribution is the reasoning-first rewriting layer, which is the missing piece.

4. ***"The knowledge base is hand-curated. Does this scale?"***

Answer: We wrote DuckDB rules in ~2 hours and achieved Xx speedup. The format is structured and DBA-extensible. Automated knowledge discovery is promising future work, but curated rules provide correctness guarantees that auto-generated rules currently cannot.

5. ***"Where does QueryTorque fail?"***

Answer: Present failure cases honestly. Queries where reasoning doesn't find an improvement, or where the rewrite is correct but not faster. Reviewers respect this and it preempts the most damaging critique.

## Execution Checklist

Concrete actions to make this positioning airtight.

| Action                               | Detail  |
|--------------------------------------|---|
| <b>Run the DuckDB experiment</b>     | R-Bot vs QueryTorque on TPC-DS 1TB, DuckDB. Prove the generalization gap is real and dramatic.  |
| <b>Highlight DSB skew queries</b>    | Focus on Query 21, 36, 78 — notorious for data skew. Show reasoning catches what cost estimators miss. Explain the skew pattern briefly to avoid cherry-picking optics. |
| <b>Benchmark on LITHE's gaps</b>     | LITHE cherry-picked "Hefty" queries. Run on the full suite. Match or beat on Hefty, dominate elsewhere. Have an answer ready for direct comparison.                     |
| <b>Measure overhead</b>              | Report wall-clock time for QueryTorque reasoning, token cost per query, number of DB round-trips. Competitors will ask. Pre-empt the question.                          |
| <b>Document the knowledge format</b> | Publish the Gap Profile structure and examples. Show how a DBA extends it. Reviewers need to see this is reproducible and not a black box.                              |
| <b>Own your failure cases</b>        | Identify queries where QueryTorque doesn't improve performance. Present them honestly. Reviewers respect this and it preempts the most damaging critique.               |

|                                   |  |
|-----------------------------------|--|
| <b>Write DuckDB rules (2 hrs)</b> | Curate a DuckDB-specific rule set to prove the portability story is practical, not theoretical. The speed of porting is itself a result. |
|-----------------------------------|--|

## Paper Sequencing

**Paper 1 (this submission):** QueryTorque with curated knowledge, reasoning-first architecture, cross-engine results on Postgres + DuckDB. Establishes the architecture and proves the approach.

**Paper 2 (follow-up):** Autonomous Knowledge Discovery for Cross-Engine Query Optimization. Full paper dedicated to proving generated knowledge is correct and useful, building on the established system.

*This is how strong research groups sequence claims. One clean contribution per paper. Don't try to solve everything at once.*

**Two pillars. Clean story. Defensible at every point.**